



# Practical security in the brave new Kubernetes world

Me harty's guide  
fer all ye scallywags



# The modern application stack



Application



Orchestrator



Container

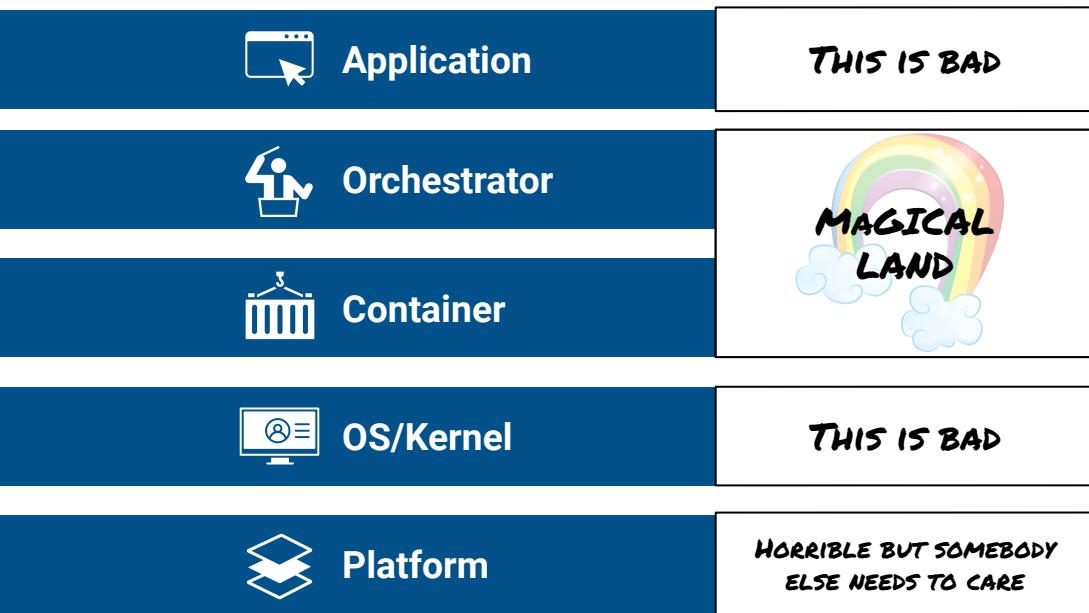


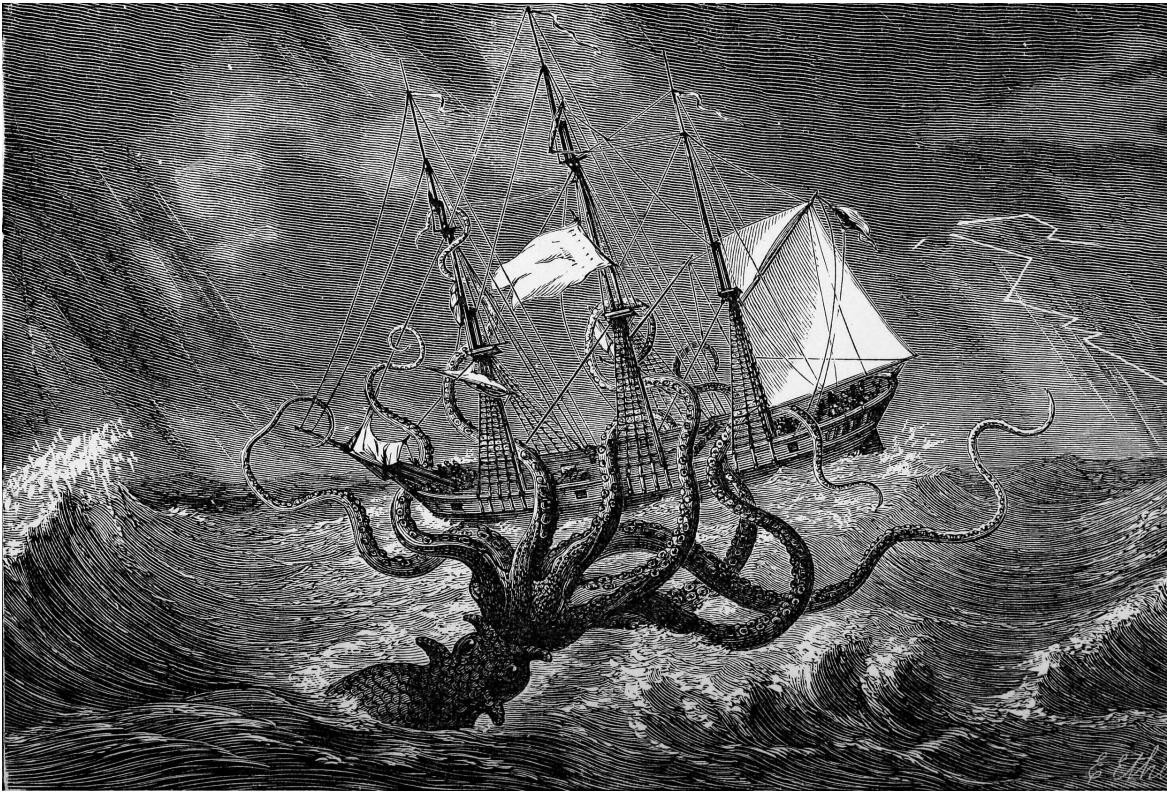
OS/Kernel



Platform

# The modern application stack







# Ahoy!

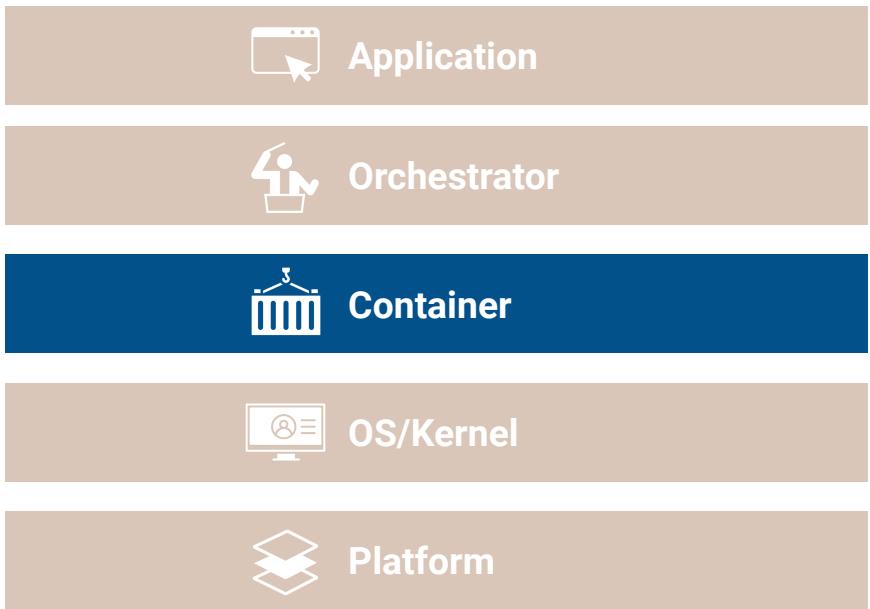


I am **Alex Ivkin**, Director of Solutions at Eclipsium, a US firmware security company, from Portland, Oregon.

I do secure deployments in (in)secure environments, including container orchestration, application security, and firmware security.



# The modern application stack



# Can containers help with my security?

Yes, if you build them, deploy and run them properly.

Types of issue a container can help with limiting the damage

- Persistence
- Tooling/living-of-the-land
- Path traversal
- Uncontrolled resource consumption/DoS



# Where containers will not help

If your app is bad, it won't fix it

- Injection/Insecure Deserialization/RCE
- Improper input validation
- Use-after-free, TOCTOU, overflow
- XSS, CSRF, SSRF



# *But wait there's more!*

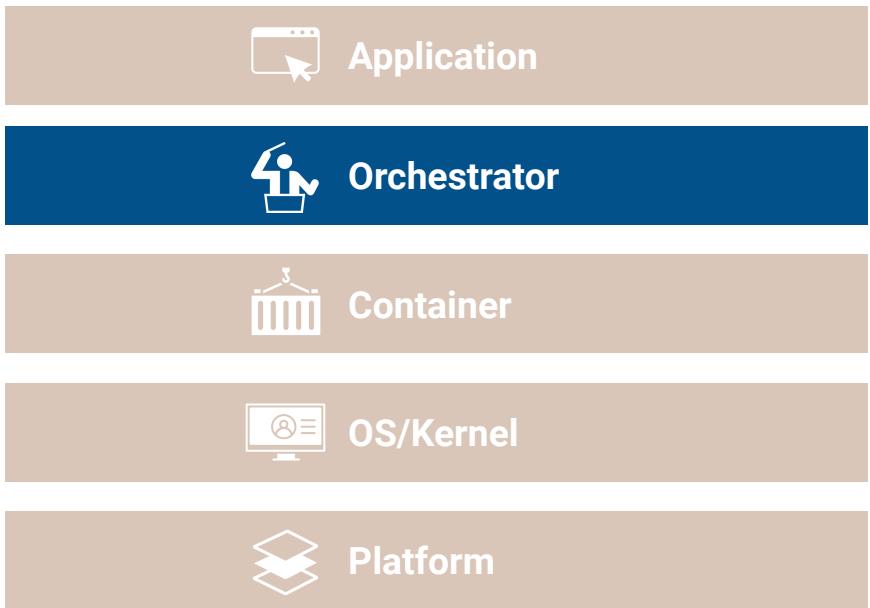
- Software supply chain, componentes with known vulnerabilities and updating these
- iptables mess
- Kernel 'exploits
- Change docker defaults, get yourself in trouble



'ow do ye ship yer containers?



# The modern application stack



# Infrastructure as a Code



Chef

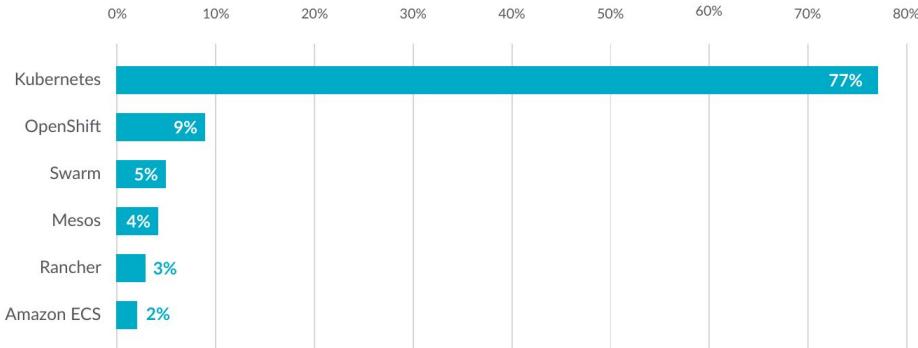


Puppet



Terraform

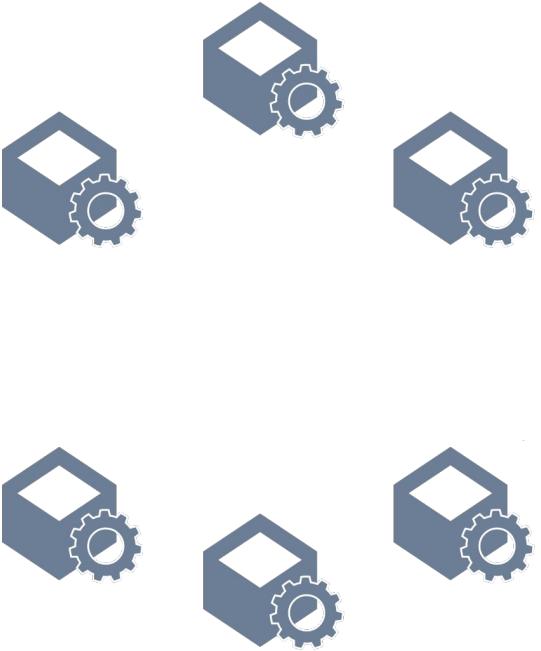
# Kubernetes!

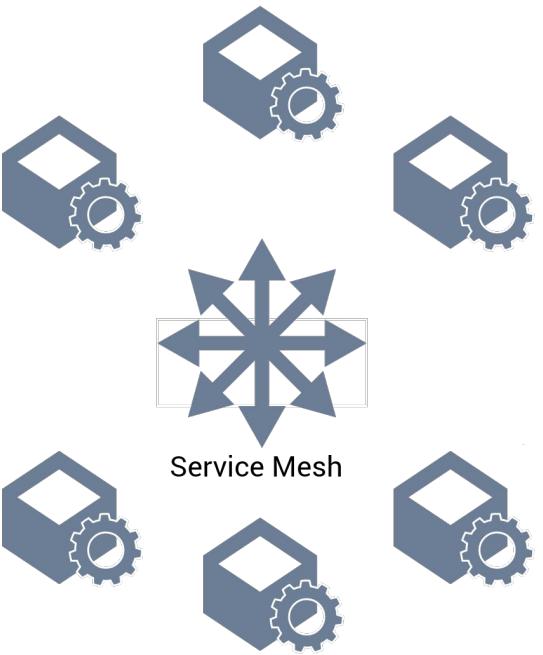


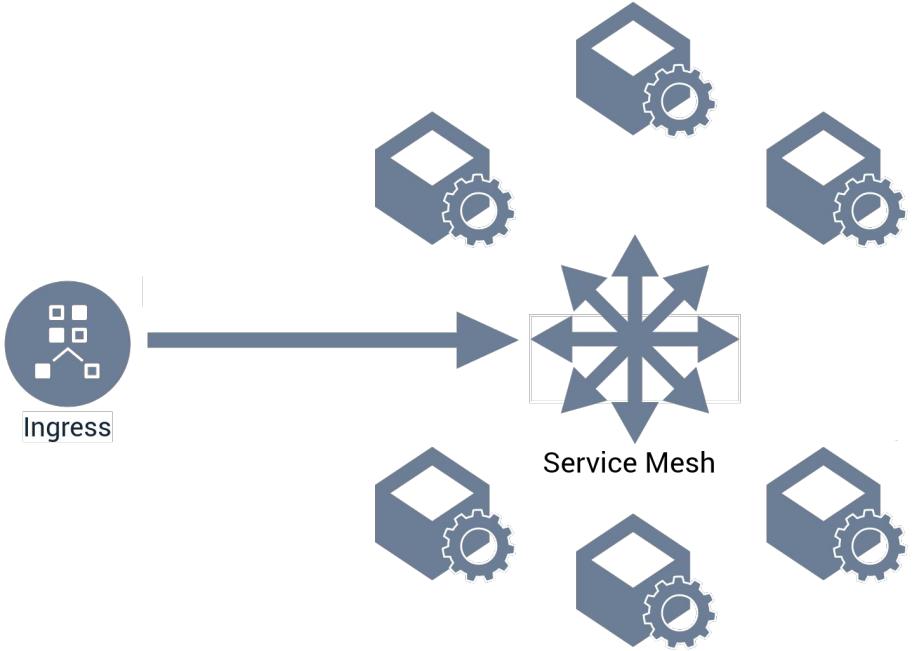
<https://sysdig.com/blog/sysdig-2019-container-usage-report/>

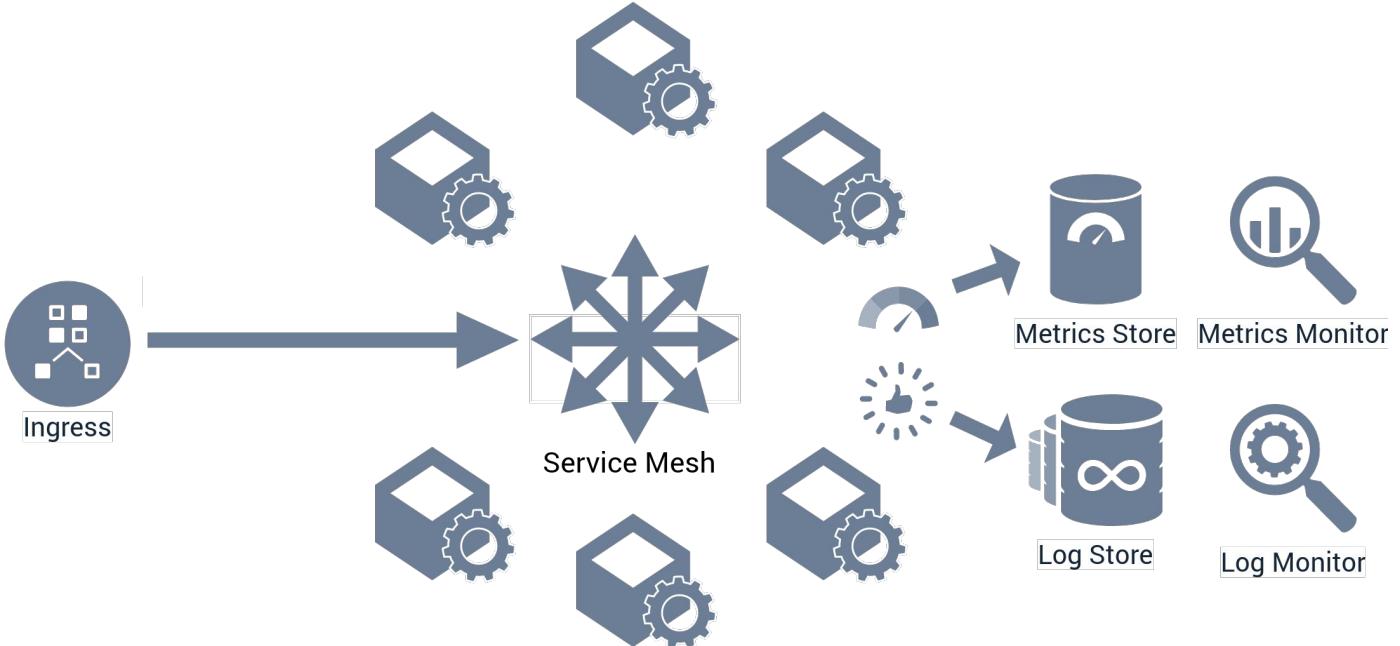
# Core K8s concepts

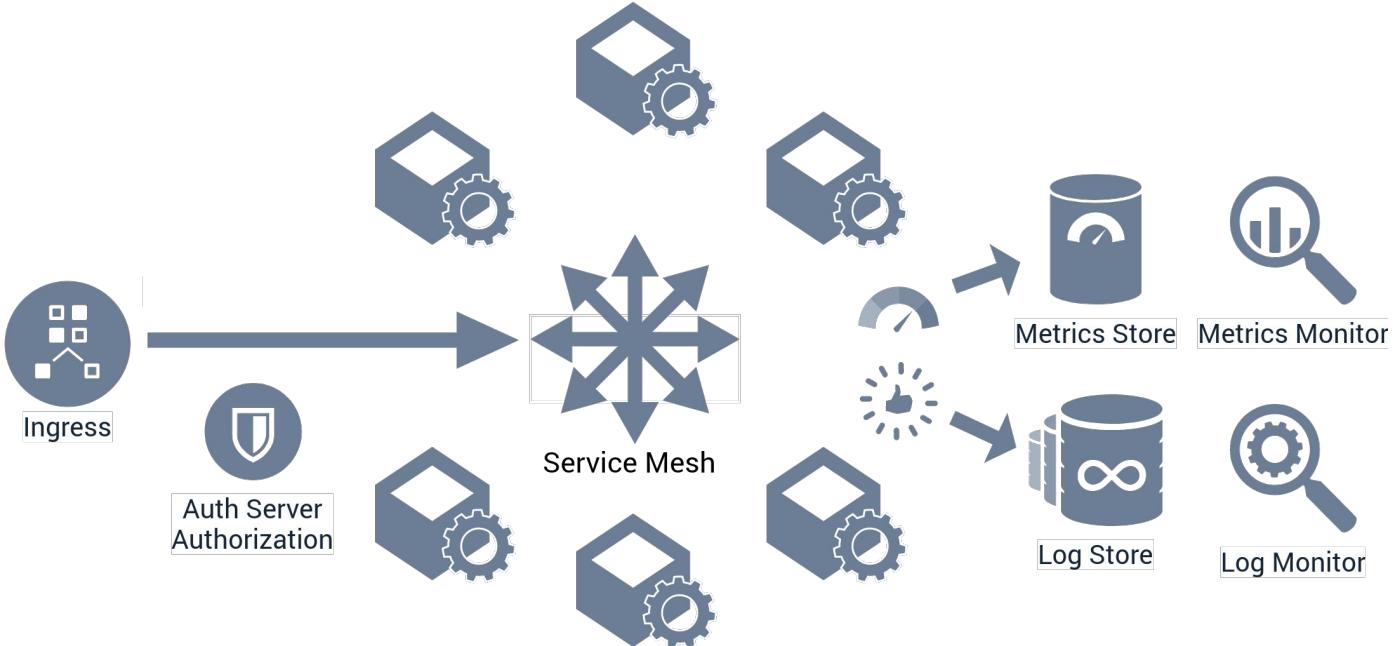
1. Declarative deployment
2. No hidden APIs
3. Meet the developer where they are at

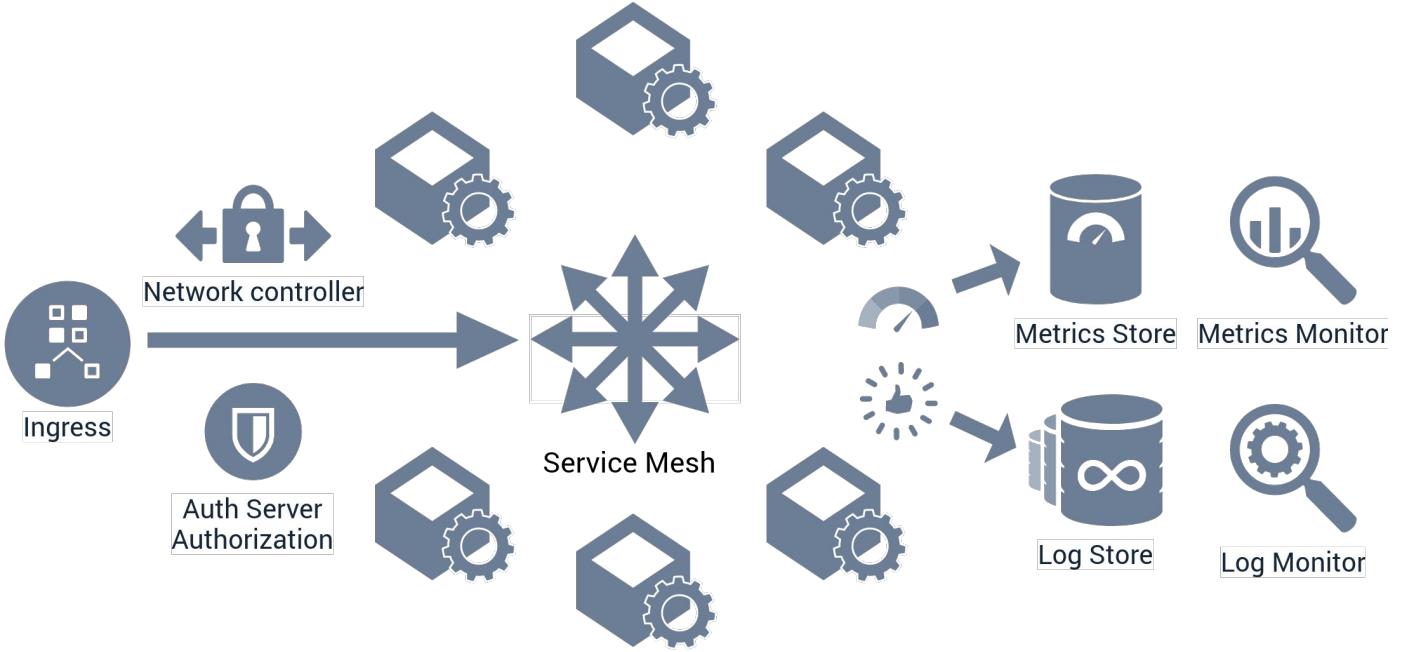


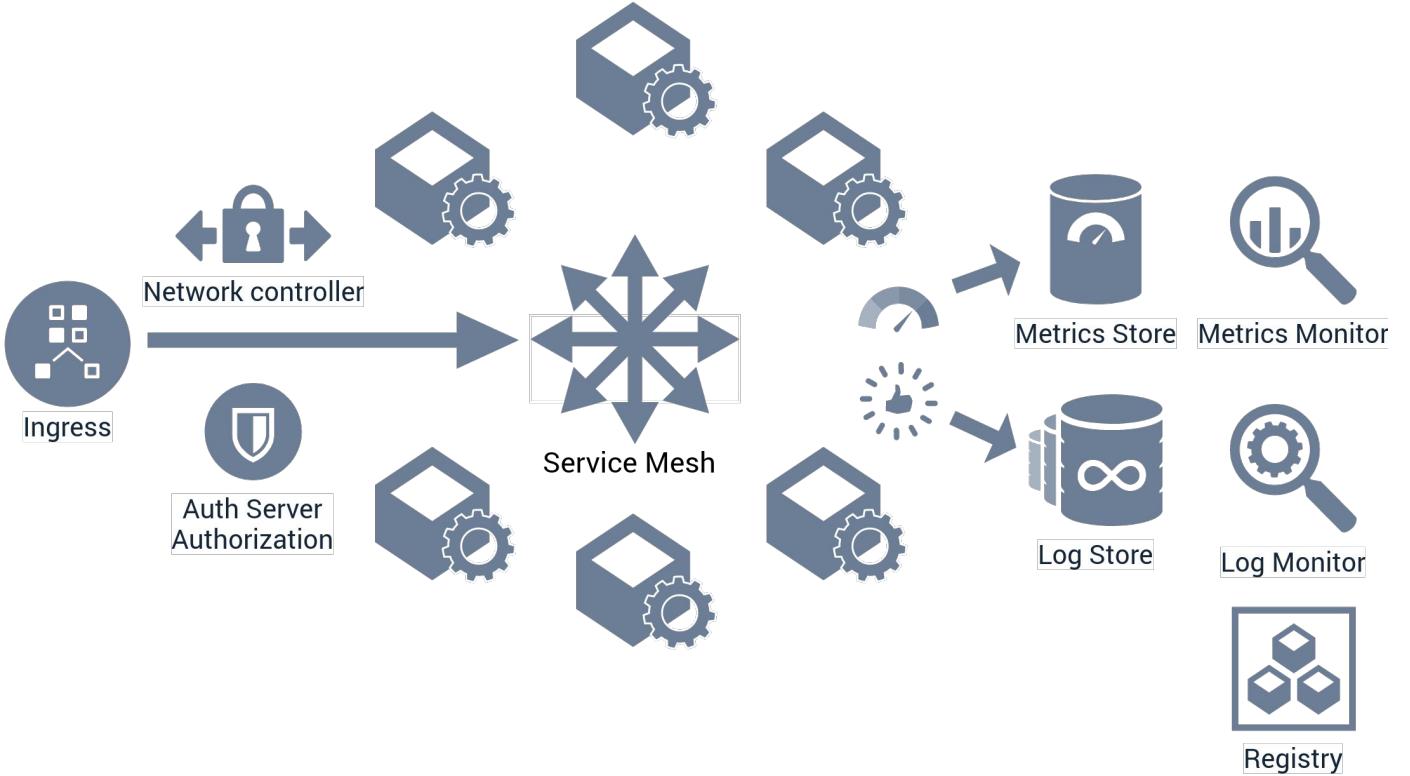


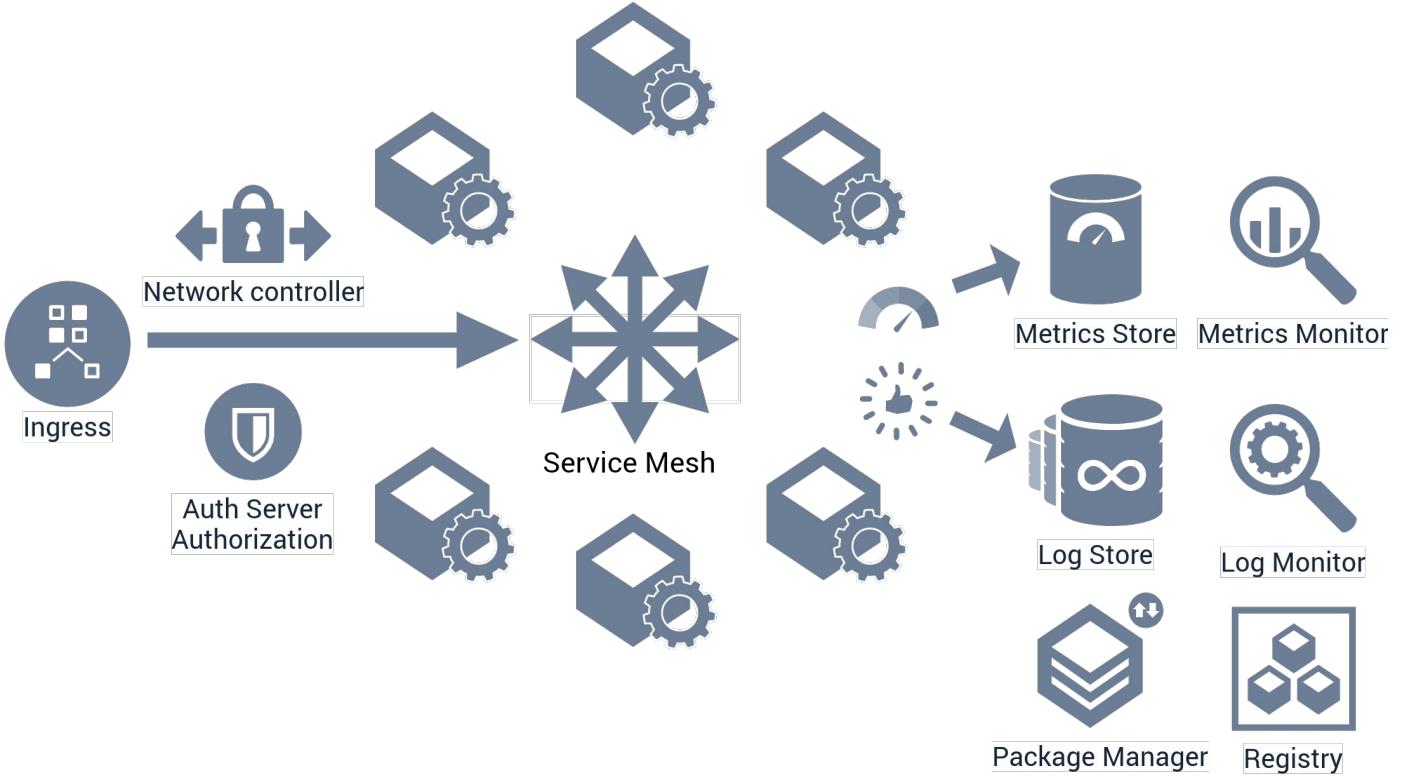














TLS Manager

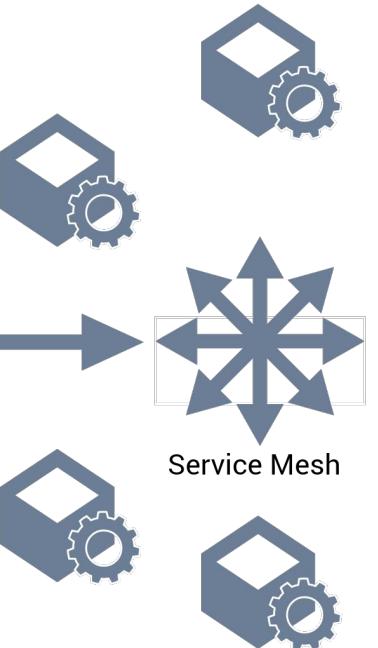


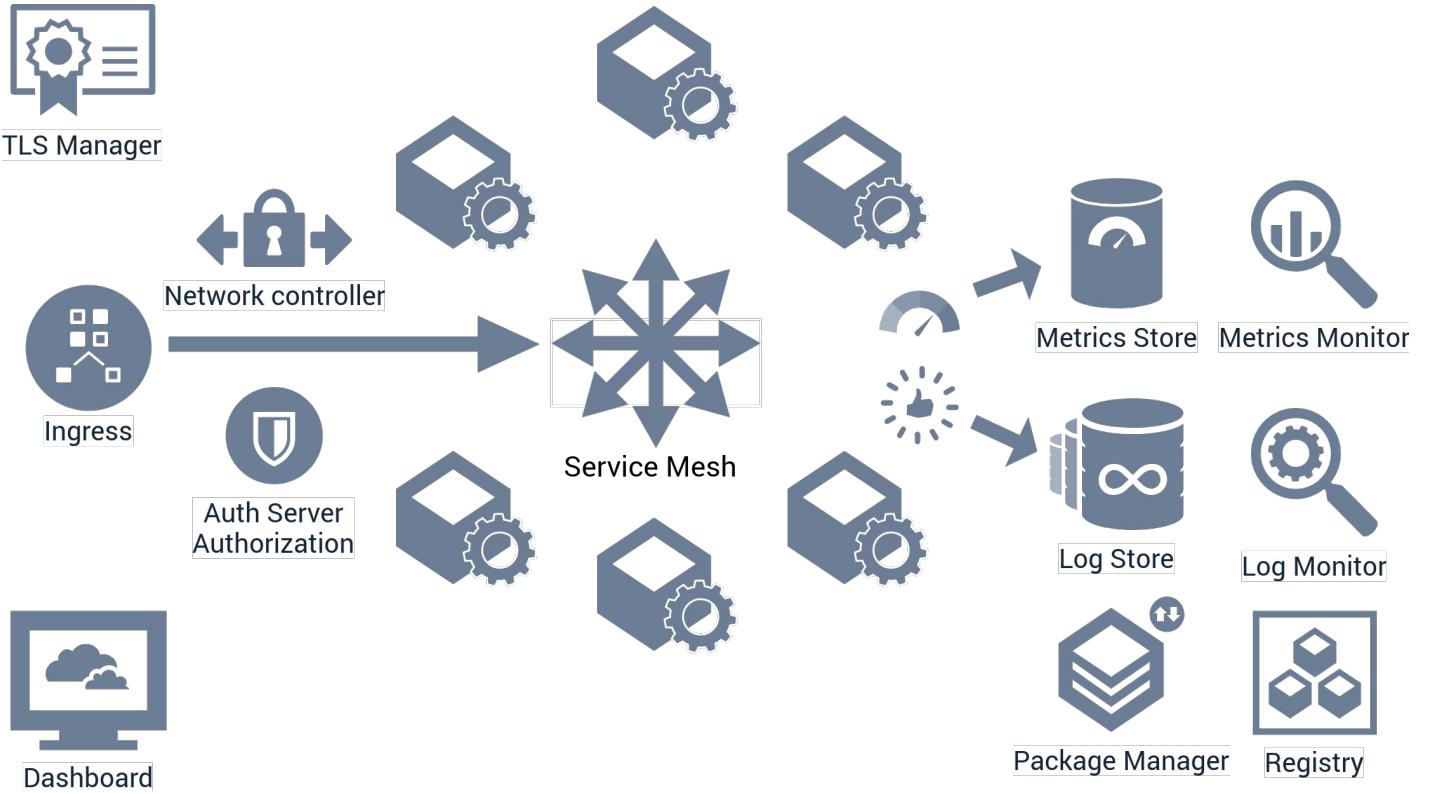
Ingress

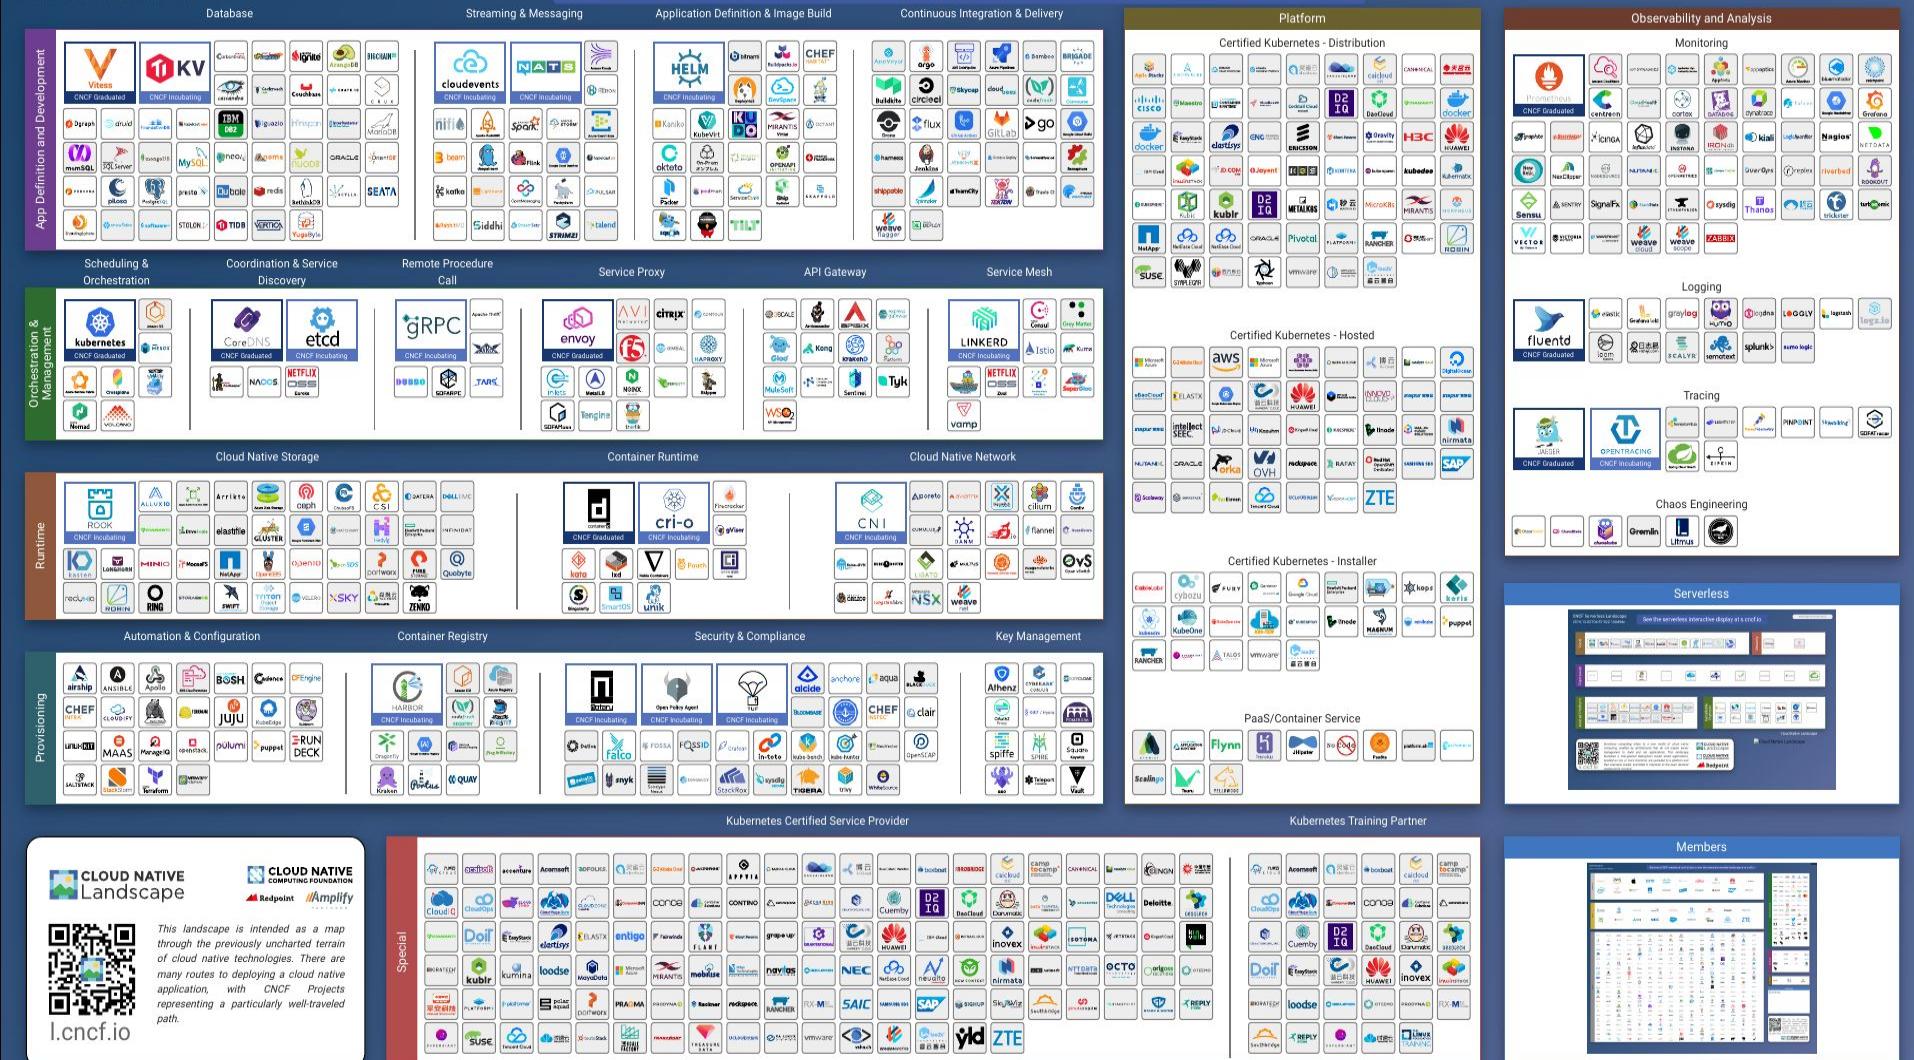
Network controller

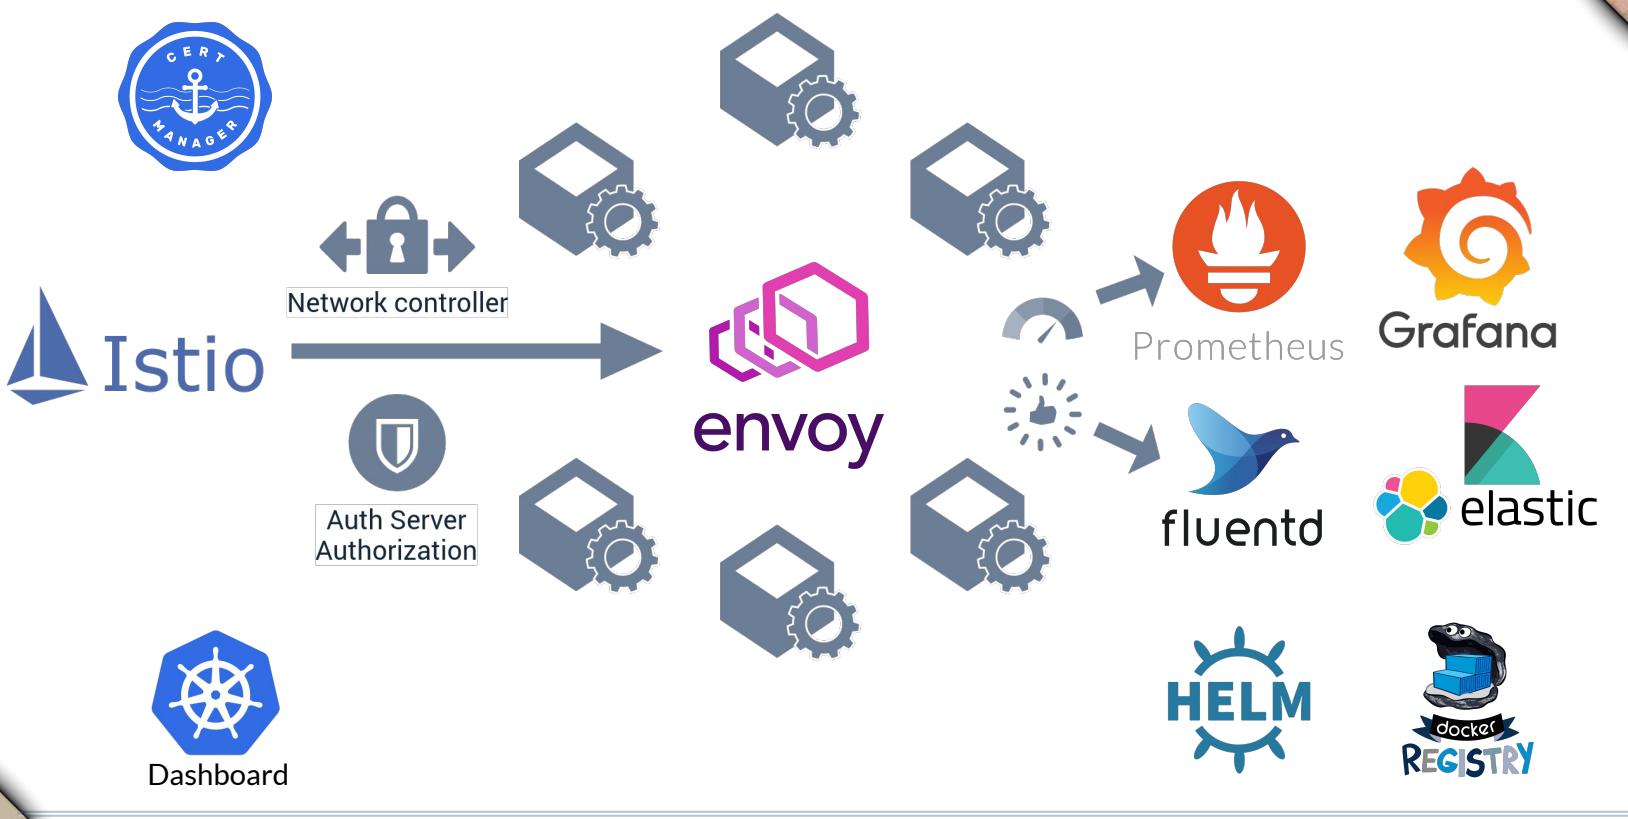


Auth Server  
Authorization











cc by 4

@alexivkinx

# Declarative deployment

- You can find everything you need in K8s resources
- Kubernetes will do what you need for you



# *No hidden APIs*

API's, API's everywhere

- K8s master API, kubelet API, Application APIs
- RBAC and network policies are the only things between you and the cluster controls.



# *Meet the developer where they are at*

- Secrets, config data and the application state are in env variables and folder mounts
- Legacy apps hastily ported to containers
- An average dev is not that good with security

# Pillage!

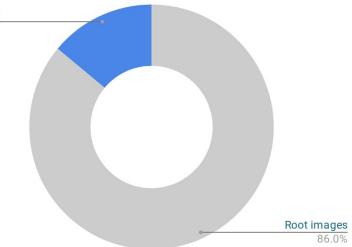
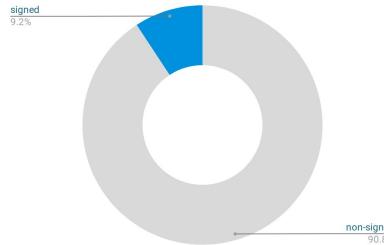


cc by 4

@alexivkinx

# Get a code exec in a pod

- Pentest
  - The app - SSRF is the best
  - The APIs
  - No need for the privesc
- Supply chain - pre-implant the image
- Look for known vulns in the k8s world



The image shows a GitHub search interface with a map of the world in the background. The search query is `is:open label:area/security`. The results are displayed in four cards:

- istio / istio**: Issues 938, Pull requests 84. Filtered to 167 Open and 791 Closed.
- grafana / grafana**: Issues 2,247, Pull requests 125. Pinned issue: **Observability Roadmap** (#15999) opened on Mar 14 by davikal, status Open.
- prometheus / prometheus**: Issues 364, Pull requests 126. Filtered to 364 Open and 2,656 Closed.
- envoyproxy / envoy**: Issues 566, Pull requests 59. Filtered to 11 Open and 21 Closed.

Each card includes a search bar and a "Clear current search query, filters, and sorts" button.

# Kubernetes Network Plugins



	Calico	Canal	Cilium	Flannel	Kube-router	WeaveNet
Encryption	No	No	Yes	No	No	Yes
Network Policies	Ingress Egress	Ingress Egress	Ingress Egress	None	Ingress	Ingress Egress

# Ingress controllers



	Ingress-nginx	Nginx-ingress	Istio	Ambassador	Traefik	Gloo
Authentication	Yes	No	mTLS, OpenID	Basic, OAuth	Basic, URL	
JWT Support	No		Yes		No	
WAF	ModSecurity		ModSecurity	No	No	

[https://docs.google.com/spreadsheets/d/1DnsHtdHbxjvHmxvlu7VhzWcWgLAn\\_Mc5L1WhLDA\\_\\_k/edit#gid=0](https://docs.google.com/spreadsheets/d/1DnsHtdHbxjvHmxvlu7VhzWcWgLAn_Mc5L1WhLDA__k/edit#gid=0)

# Real world Kubernetes

- Broken Authentication, Access Control
- Security Misconfiguration
- Insufficient Logging and Monitoring
- Service accounts/secrets distributions
- mTLS and In-transit encryption



# What do I do?



# Easy

- Use simple images to start from or grab and reconstruct dockerfile and build the image yourself
- No privileged containers or shared volumes with the node
- Monitor for rogue containers
- In the cloud - secure metadata, use COS/Bottlerocket
- On baremetal - use hardened OS (Flatcar, Minimal Debian, Alpine with CRI)
- Use RBAC (duh!)



<https://github.com/alexivkin/docker-historian>

# Normal

- Build from scratch or use distroless images
- Use PodSecurity Policies
- Have multiple registries and authenticate/authorize access
- Remap root to non-root users for all containers
- Upgrade master nodes and all nodes to the latest version
- SAST/DAST/SCA for images
- Ship out K8s secrets to a vault



# Hardmode

- Use Pod Admission Policies
- Aim for zero-trust (mTLS, SPIFFE/SPIRE)
- Sign your images
  - ➔ Use TUF/Notary to address supply chain risks
- Unmix sensitive workloads
  - ➔ Container separation is weaker than hypervisors
  - ➔ Namespaces will not provide hard multitenancy
    - ➔ Use separate clusters



# Security best practices, frameworks and standards

- CTFs and playgrounds
- CIS Benchmarks for Docker and Kubernetes
- Aqua, StackRox, SysDig, Tigera, Neuvector whitepapers
- Night time reading
  - US DoD DevSecOps reference design
  - NIST SP 800-190 - Compliance in Container Environments
  - SANS Checklist for Audit of Docker containers
- <https://github.com/alexivkin/containerpwn>
- <https://github.com/alexivkin/kubepwn>



**It's gonna get more complex  
Errors, not vulnerabilities, lead to a lot of breaches  
Make it secure by default**



<https://github.com/alexivkin> 

@alexivkinx 

in/alexivkin 

